

The Open Discussion (The Technical Veto)

■■ Friday, Sept 25, 2026: The Open Discussion (The Technical Veto)

The hardest job in engineering leadership right now isn't learning how to deploy new models, it's having the technical courage to say **"no"** to them.

While we often point to broader market trends or product strategy as the driving force behind the rush to adopt AI, the reality is that the push can sometimes originate from within our own engineering teams as well.

Lately, there is a frustrating trend where developers push to implement complex Generative AI pipelines for problems that require basic data science or pure, deterministic code.

Why? Because it's trendy. Because no one wants to write a standard SQL query or an optimized matching filter when they can put "Built an Agentic LLM framework" on their resume.

They rush to slap an API endpoint or a heavy vector stack onto a basic sorting task without understanding the underlying mechanics:

- They don't calculate the token unit economics.
- They ignore the massive user-facing latency taxes.
- They fail to account for the risk of probabilistic models lying to customers.

Our primary responsibility as technical leaders is to protect our core systems from this hype cycle, even when it comes from our own teams. We must force our developers to anchor themselves in engineering pragmatism, not resume-driven development.

To close out our first week, I'd love to open the floor to this community:

How do you handle developers/anyone on your team who insist on using complex LLM frameworks for tasks that simple, traditional code could solve in milliseconds? Drop your number below:

- 1■■ Strict architectural veto (Enforce simple code/math)
- 2■■ Let them prototype it first, then let the latency/cost metrics kill it
- 3■■ Struggle to push back due to organizational AI pressure

Let's talk in the comments. ■